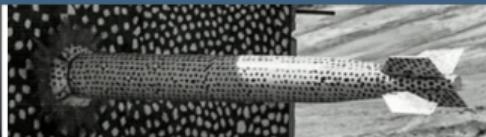




Sandia  
National  
Laboratories

# PyTrilinos2: automatic (re)generation of a Python interface for Trilinos



*Presented by:*

K. Liegeois & C. Glusa

Trilinos User Group, October 2023



Sandia National Laboratories is a multimission laboratory managed and operated by National Technology & Engineering Solutions of Sandia, LLC, a wholly owned subsidiary of Honeywell International Inc., for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NAD003525.

## Motivations for a Python interface for Trilinos



Python is a glue language that allows to:

- ▶ ease the prototyping of new methods and applications,
- ▶ improve the pre and post processing workflow,
- ▶ drive complex software through user-friendly front-ends,
- ▶ manage inter-operation between libraries and software written in different languages.

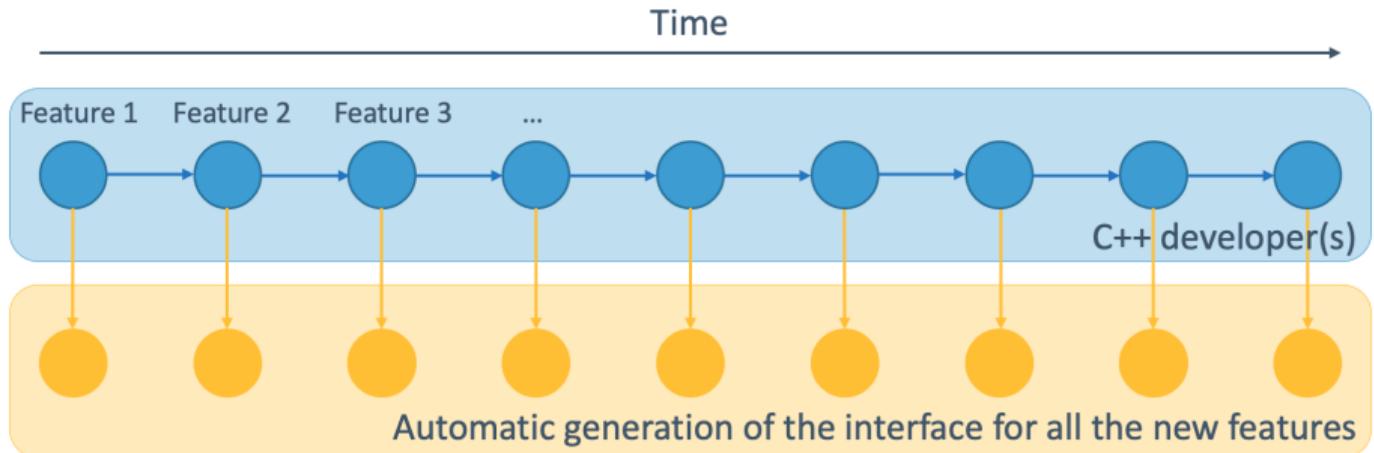
## Motivations for a Python interface for Trilinos



In particular, a Python interface for Trilinos could allow to:

- ▶ drive parametric computations from Python:
  - ▶ perform uncertainty quantification analysis,
  - ▶ optimize parameter values,
  - ▶ compute sensitivity w.r.t parameters using finite differences, ...
- ▶ prototype new solvers while relying on existing C++ features:
  - ▶ new multigrid cycles for MueLu,
  - ▶ new mixed precision linear solvers for Belos,
  - ▶ new solvers for Ifpack2, ...
- ▶ embed Machine Learning (ML) models written in Python into existing C++ code:
  - ▶ use a ML model at an integration point to evaluate as a constitutive model while assembling matrices, ...
- ▶ lower the learning curve associated to the usage of Trilinos and potentially increase the user base.

## PyTrilinos2: goals of the reboot



### Goals:

- ▶ Automatic generation of the interface for new packages.
- ▶ Automatic update of the interface to reflect C++ changes.
- ▶ No latency between new C++ features and their Python interface.

## Approach used for PyTrilinos2



PyTrilinos2 relies on the combination of two tools: Pybind11 and Binder.

### **Pybind11:**

- ▶ a lightweight header-only library that exposes C++ types in Python and vice versa,
- ▶ similar to Boost.Python,
- ▶ supports custom smart pointers (such as Teuchos::RCP), template arguments, inheritance, ...
- ▶ can be easily installed using pip or spack.

### **Binder:**

- ▶ tool for the automatic generation of Python bindings using Pybind11 and Clang LibTooling libraries,
- ▶ generates the Pybind11 lines in .cpp files but does not compile them,
- ▶ developed at the Johns Hopkins University by Sergey Lyskov,
- ▶ can be installed using spack.

# Example of how to add an extra package in the interface



```
#include <PyTrilinos2_Teuchos_ETI.hpp>
#include <PyTrilinos2_Tpetra_ETI.hpp>
#include <PyTrilinos2_MueLu_ETI.hpp> // <----- One new header for the new package
```

```
#ifndef PYTRILINOS2_MUELU_ETI
#define PYTRILINOS2_MUELU_ETI
#include <MueLu_CreateTpetraPreconditioner.hpp>

#define BINDER_MUELU_CREATETPETRAPRECONDITIONER_INSTANT(SCALAR, LO, GO, NO) \
template Teuchos::RCP<MueLu::TpetraOperator<SCALAR, LO, GO, NO> > \
CreateTpetraPreconditioner<SCALAR, LO, GO, NO> \
const Teuchos::RCP<Tpetra::Operator<SCALAR, LO, GO, NO> > &inA, \
Teuchos::ParameterList& inParamList); \
\\

namespace MueLu {
BINDER_MUELU_CREATETPETRAPRECONDITIONER_INSTANT(
    Tpetra::Details::DefaultTypes::scalar_type,
    Tpetra::Details::DefaultTypes::local_ordinal_type,
    Tpetra::Details::DefaultTypes::global_ordinal_type,
    Tpetra::KokkosClassic::DefaultNode::DefaultNodeType)
}
#endif // PYTRILINOS2_MUELU_ETI
```

## Extra work needed for adding a new package



The binder configuration script needs to be updated to not bind what should not be bound:

```
-class MueLu::FactoryAcceptor  
-class MueLu::FactoryFactory  
-class MueLu::FactoryManagerBase  
-class MueLu::FactoryManager  
-class MueLu::FactoryBase  
-class MueLu::Hierarchy  
#...
```

The PyTrilinos2 CMake files need to be updated to add the new package:

```
....  
list(APPEND BINDER_OPTIONS --bind Teuchos)  
list(APPEND BINDER_OPTIONS --bind Tpetra)  
list(APPEND BINDER_OPTIONS --bind MueLu) # <---- New line  
#...
```

```
....  
SET(LIB_REQUIRED_DEP_PACKAGES  
    Teuchos  
    Tpetra  
    MueLu # <---- New line  
)  
#...
```

## Example of **automatically** generated bindings for the new package

```
void bind_PyTrilinos2_50(std::function< pybind11::module &(std::string const &namespace_) > &M)
{
{ // MueLu::SingleLevelFactoryBase file:MueLu_SingleLevelFactoryBase.hpp line:64
    pybind11::class_<MueLu::SingleLevelFactoryBase, Teuchos::RCP<MueLu::SingleLevelFactoryBase>, MueLu::Factory>
        cl(M("MueLu"), "SingleLevelFactoryBase", "Base class for factories that use one level (currentLevel).");
}
{ // MueLu::TwoLevelFactoryBase file:MueLu_TwoLevelFactoryBase.hpp line:67
    pybind11::class_<MueLu::TwoLevelFactoryBase, Teuchos::RCP<MueLu::TwoLevelFactoryBase>, MueLu::Factory>
        cl(M("MueLu"), "TwoLevelFactoryBase",
            "Base class for factories that use two levels (fineLevel and coarseLevel).");
}
{ // MueLu::SetFactoryManager file:MueLu_HierarchyUtils_decl.hpp line:65
    pybind11::class_<MueLu::SetFactoryManager, Teuchos::RCP<MueLu::SetFactoryManager>>
        cl(M("MueLu"), "SetFactoryManager", "An exception safe way to call the method 'Level::SetFactoryManager()' ");
        cl.def( pybind11::init( [] (MueLu::SetFactoryManager const &o){ return new MueLu::SetFactoryManager(o); } ) );
}
// MueLu::CreateTpetraPreconditioner(const class Teuchos::RCP<class ...
M("MueLu").def("CreateTpetraPreconditioner", (class Teuchos::RCP<class MueLu::TpetraOperator<double, int, long long,
    class Tpetra::KokkosCompat::KokkosDeviceWrapperNode<class Kokkos::OpenMP, class Kokkos::HostSpace> > >
    (*) (const class Teuchos::RCP<class Tpetra::Operator<double, int, long long,
        class Tpetra::KokkosCompat::KokkosDeviceWrapperNode<class Kokkos::OpenMP, class Kokkos::HostSpace> > > &,
        class Teuchos::ParameterList &) ) &MueLu::CreateTpetraPreconditioner<double,int,long long>,
    Tpetra::KokkosCompat::KokkosDeviceWrapperNode<Kokkos::OpenMP, Kokkos::HostSpace>>,
    "Helper function to create a MueLu or AMGX preconditioner that can be used by Tpetra...", 
    pybind11::arg("inA"), pybind11::arg("inParamList"));
}
```

## Example of usage of PyTrilinos2



The presented example uses Teuchos, Tpetra, and MueLu.

Why have we started with Tpetra? Tpetra:

- ▶ relies heavily on templates,
- ▶ uses MPI communicator,
- ▶ is the base package of a stack of solvers,
- ▶ relies on Kokkos for shared memory parallelism,
- ▶ uses ETI and allows to instantiate multiple instances of one class within one build.

The current work branch of PyTrilinos2 provides binding for MultiVector, Vector, Map, CrsGraph, CrsMatrix, Export, Import, ...

# Example: 1D Laplacian solved with CG



```
def main():
    comm = Teuchos.getTeuchosComm(MPI.COMM_WORLD)

    n = 3000
    A = assemble1DLaplacian(n, comm)
    mapT = A.getRowMap()

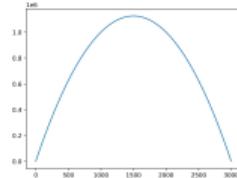
    x = get TypeName('Vector')(mapT, True)
    b = get TypeName('Vector')(mapT, False)
    b.putScalar(1.)
    pl = Teuchos.ParameterList()
    P = MueLu.CreateTpetraPreconditioner(A, pl)

    CG(A, x, b, P, max_iter=30)
    n0 = 0
    if comm.getRank() == 0:
        n0 = n
    mapT0 = type(mapT)(n, n0, 0, comm)

    x0 = get TypeName('Vector')(mapT0, True)
    export = get TypeName('Export')(mapT0, mapT)
    x0.doImport(source=x, exporter=export,
                CM=Tpetra.CombineMode.REPLACE)

    if comm.getRank() == 0:
        plt.figure()
        plt.plot(x0.getLocalViewHost())
        plt.savefig('x0_view.png')
```

- ▶ Solve a 1D Laplacian using CG,
- ▶ the CG is written in Python using Tpetra function calls,
- ▶ the script can be run with MPI,
- ▶ the postprocess is done on the rank 0 after importing the solution vector and getting its local view host (as a NumPy array),
- ▶ Kokkos is used under the hood and works with CUDA backend.



# PyROL

The same strategy has been apply to ROL in collaboration with [Aurya Javeed](#).

Available features:

- ▶ The python interface exposes the ROL virtual classes,
- ▶ The user can define a child class in Python to define their own vector type.
- ▶ The user can define their objective functions in Python.
- ▶ The user can then call the optimizer from Python.

PyROL will soon be available and user will be able to install it from source (as PyTrilinos2) or using pip.

```
# Matrix from rol/example/quadratic/example_01.cpp
class matrix(LinearOperator):
    def __init__(self, dim):
        self.dim = dim
        super().__init__()
    def apply(self, Hv, v, tol):
        for i in range(0, self.dim):
            Hv[i] = 2.*v[i]
            if i > 0:
                Hv[i] -= v[i-1]
            if i < self.dim - 1:
                Hv[i] -= v[i+1]

op = matrix(10)
g = vector_type.full(10, 0.)
x = vector_type.full(10, 1.)
zero = vector_type.full(10, 0.)

bnd = Bounds_double_t(zero, x.clone())
obj = QuadraticObjective(op, g)

params = getParametersFromXmlFile("input.xml")
status = StatusTest(params)

problem = Problem_double_t(obj, x)
problem.addBoundConstraint(bnd)
solver = Solver_double_t(problem, params)
cout = openOfstream('test.txt')
solver.solve(cout, status)
closeOfstream(cout)

state = solver.getAlgorithmState()
```

## Current stage and future work



Current stage:

- ▶ Binder can be installed using spack.
- ▶ The approach has been applied on ROL.
- ▶ Initial PyTrilinos2 PR is open for review.

Future works:

- ▶ The content of the PR will be tested in a nightly build.
- ▶ PyTrilinos(1) will be deprecated.
- ▶ Create the interface for other packages/features.
- ▶ Improve PyTrilinos2 CMake logic to allow optional packages.
- ▶ Potentially couple the work with PyKokkos.

## Example: Tpetra CG



```
def CG(A, x, b, prec, max_iter=20, tol=1e-8):
    r = type(b)(b, Teuchos.DataAccess.Copy)
    A.apply(x,r,Teuchos.ETransp.NO_TRANS,alpha=-1,beta=1)

    p = type(r)(r, Teuchos.DataAccess.Copy)
    q = type(r)(r, Teuchos.DataAccess.Copy)

    Br = type(r)(r, Teuchos.DataAccess.Copy)
    prec.apply(r, p)
    gamma = sqrt(r.dot(p))

    if gamma < tol:
        return 0
    for j in range(max_iter):
        A.apply(p, q)
        c = q.dot(p)
        alpha = gamma**2 / c
        x.update(alpha, p, 1)
        r.update(-alpha, q, 1)
        prec.apply(r, Br)
        gamma_next = sqrt(Br.dot(r))
        beta = gamma_next**2/gamma**2
        gamma = gamma_next
        if gamma < tol:
            return j+1
        p.update(1, Br, beta)
    return max_iter
```

# Example: 1D Laplacian



```
def assemble1DLaplacian(n, comm):
    mapT = get TypeName('Map')(n, 0, comm)
    graph = get TypeName('CrsGraph')(mapT, 3)
    for i in range(mapT.getMinLocalIndex(), mapT.getMaxLocalIndex() + 1):
        global_i = mapT.getGlobalElement(i)
        indices = [global_i]
        if global_i > 0:
            indices.append(global_i - 1)
        if global_i < mapT.getMaxAllGlobalIndex():
            indices.append(global_i + 1)
        graph.insertGlobalIndices(global_i, indices)
    graph.fillComplete()

    A = get TypeName('CrsMatrix')(graph)
    for i in range(mapT.getMinLocalIndex(), mapT.getMaxLocalIndex() + 1):
        global_i = mapT.getGlobalElement(i)
        indices = [global_i]
        vals = [2.]
        if global_i > 0:
            indices.append(global_i - 1)
            vals.append(-1.)
        if global_i < mapT.getMaxAllGlobalIndex():
            indices.append(global_i + 1)
            vals.append(-1.)
        A.replaceGlobalValues(global_i, indices, vals)
    A.fillComplete()
    return A
```