




# A Distributed-Memory Schur-complement PCA Preconditioner for Gemma Ill-conditioned Problems



Vinh Dang and Joseph Kotulski (1352)

[vqdang@sandia.gov](mailto:vqdang@sandia.gov), [jkotul@sandia.gov](mailto:jkotul@sandia.gov)

Trilinos User Group Meeting 2023

October 30th - November 2nd




Sandia National Laboratories is a multimission laboratory managed and operated by National Technology & Engineering Solutions of Sandia, LLC, a wholly owned subsidiary of Honeywell International Inc., for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA0003525.

SAND 1727220

## Outline



- Overview of Electromagnetic Radiation (EMR) Problem
- Motivation and Objective
- Performance Portability
- Overview of the Schur-complement Principal Component Analysis (PCA) Preconditioner
- Parallel Implementation
  - Distributed Binary Tree
  - Schur-complement PCA Factorization
  - Generalized Conjugate Residual (GCR) Solver
- Numerical Results
- Conclusions and Future Work




# Electromagnetic Radiation (EMR) Problem Overview


## EMR Problem Overview




Electromagnetic (EM) energy couples into systems in many different ways. Our focus is energy coupling through mechanical seams or joints in the system housing, which acts as a good but imperfect EM shield. The joints form slots that allow EM energy into the system. Therefore, the problem has two parts that must be well-characterized: the slot and the cavity.



## Shielding Effectiveness against Frequency for the Higgins Cylinder




The 2" slot is halfway up the cylinder. A monopole probe is located inside the cylinder on one of the ends



Cylinder shielding effectiveness  $SE = 20 \log_{10} \left( \frac{E_{\text{interior}}}{E_{\text{exterior}}} \right)$   
in the mode-stirred and anechoic chambers


[1] Matthew B. Higgins and Dawna R. Charley, *Electromagnetic Radiation (EMR) Coupling to Complex Systems: Aperture Coupling into Canonical Cavities in Reverberant and Anechoic Environments and Model Validation*, SAND2007-7931

## Shielding Effectiveness against Frequency for the Higgins Cylinder



The 2" slot is halfway up the cylinder. A monopole probe is located inside the cylinder on one of the ends

- When the slot is at a resonance, it provides a larger drive to the cavity  
→ the cavities modes have larger field magnitudes
- Therefore, the slot and the cavity have to be well-characterized *at resonances*



[1] Matthew B. Higgins and Dawna R. Charley, *Electromagnetic Radiation (EMR) Coupling to Complex Systems: Aperture Coupling into Canonical Cavities in Reverberant and Anechoic Environments and Model Validation*, SAND2007-7931



# Motivation and Objective

## SNL's Electromagnetic Code Gemma

- Frequency-domain EM
- Hybrid approach, using:
  - A narrow slot sub-cell model to represent slots and the corresponding coupling
  - Surface integral equation method for outer region and inner region of cavity

Electric field integral equation (EFIE), for simplification:


$$L\{\mathbf{J}_S\} = \frac{1}{j\omega\mu} \hat{\mathbf{n}} \times \mathbf{E}_{\text{inc}}$$

Expand unknown in a set of RWG basis functions:

$$\mathbf{J}_S(\mathbf{r}) \approx \sum_n I_n \mathbf{f}_n(\mathbf{r})$$

Test integral equation with RWG basis functions:

$$\int_S \mathbf{f}_m \cdot L\{\mathbf{J}_S\} ds = \frac{1}{j\omega\mu} \int_S \mathbf{f}_m \cdot (\hat{\mathbf{n}} \times \mathbf{E}_{\text{inc}}) ds$$




- When an iterative method is used to solve the large, dense, complex matrix equation system:
  - $\mathbf{A}$  is extremely ill-conditioned for the problem of interest
  - Effective and efficient preconditioning is required

## Objective




- A preconditioner using Schur-complement and Principal Component Analysis (PCA) on distributed-memory accelerator-based computing platforms for dense matrices (an extension of our collaboration with Ohio State University EM group)
  - Propose a hierarchical binary distribution scheme for the binary tree
  - Target performance portability
  - For use with the Multilevel Fast Multipole Method (MLFMM) or the Adaptive Cross Approximation (ACA) for future work



# Performance Portability

# Performance Portability with Kokkos and Kokkos Kernels




- Gemma is written on top of Kokkos (an open-source **productive, portable, performant**, shared-memory programming model) and Kokkos Kernels (a library for node-level, performance-portable, computational kernels for sparse/dense linear algebra and graph operations)

[2] C. Trott *et al.*, "The Kokkos EcoSystem: Comprehensive Performance Portability for High Performance Computing," in *Computing in Science & Engineering*, vol. 23, no. 5, pp. 10-18, 1 Sept.-Oct. 2021



# Overview of the Schur-complement PCA Preconditioner

## Schur-complement PCA Preconditioner



$$(A_1^l)^{-1} = \begin{bmatrix} I & 0 \\ -(A_2^{l+1})^{-1}C_{21}^{l+1} & I \end{bmatrix} \begin{bmatrix} (I - (A_1^{l+1})^{-1}C_{12}^{l+1}(A_2^{l+1})^{-1}C_{21}^{l+1})^{-1} & 0 \\ 0 & I \end{bmatrix}$$

$$\begin{bmatrix} I & -(A_1^{l+1})^{-1}C_{12}^{l+1} \\ 0 & I \end{bmatrix} \begin{bmatrix} (A_1^{l+1})^{-1} & 0 \\ 0 & (A_2^{l+1})^{-1} \end{bmatrix}$$

Factorization using PCA:


$$(A_1^{l+1})^{-1}C_{12}^{l+1} \xrightarrow{PCA} L_1^{l+1}D_1^{l+1}R_1^{l+1} = L_1^{l+1}R_1^{l+1}$$

$$(A_2^{l+1})^{-1}C_{21}^{l+1} \xrightarrow{PCA} L_2^{l+1}D_2^{l+1}R_2^{l+1} = L_2^{l+1}R_2^{l+1}$$

Inverse of Schur-complement:


$$(I - (A_1^{l+1})^{-1}C_{12}^{l+1}(A_2^{l+1})^{-1}C_{21}^{l+1})^{-1} \approx (I - L_1^{l+1}R_1^{l+1}L_2^{l+1}R_2^{l+1})^{-1}$$

$$= I + L_1^{l+1}(I - R_1^{l+1}L_2^{l+1}R_2^{l+1}L_1^{l+1})^{-1}R_1^{l+1}L_2^{l+1}R_2^{l+1} = I + L_1^{l+1}S^lR_2^{l+1}$$



**Store:**  $L_1^{l+1}(m \times k), R_1^{l+1}(k \times n)$   
 $L_2^{l+1}(n \times q), R_2^{l+1}(q \times m)$   
 $S^l(k \times q)$

## Schur-complement PCA Preconditioner



Level  $l-1$   
 $A_1^{l-1}$   
 Level  $l$

$$(A_1^{l-1})^{-1} = \begin{bmatrix} I & 0 \\ -(A_2^l)^{-1}C_{21}^l & I \end{bmatrix} \begin{bmatrix} (I - (A_1^l)^{-1}C_{12}^l(A_2^l)^{-1}C_{21}^l)^{-1} & 0 \\ 0 & I \end{bmatrix}$$

$$\begin{bmatrix} I & -(A_1^l)^{-1}C_{12}^l \\ 0 & I \end{bmatrix} \begin{bmatrix} (A_1^l)^{-1} & 0 \\ 0 & (A_2^l)^{-1} \end{bmatrix}$$

Factorization using PCA:


$$(A_1^l)^{-1}C_{12}^l \xrightarrow{PCA} L_1^l D_1^l R_1^l = L_1^l R_1^l$$

$$(A_2^l)^{-1}C_{21}^l \xrightarrow{PCA} L_2^{l+1} D_2^{l+1} R_2^{l+1} = L_2^{l+1} R_2^{l+1}$$

Inverse of Schur-complement:

$$(I - (A_1^l)^{-1}C_{12}^l(A_2^l)^{-1}C_{21}^l)^{-1} \approx (I - L_1^l R_1^l L_2^l R_2^l)^{-1}$$

$$= I + L_1^l (I - R_1^l L_2^l R_2^l L_1^l)^{-1} R_1^l L_2^l R_2^l = I + L_1^l S^{l-1} R_2^l$$




**Store:**  $L_1^l (m' \times k')$ ,  $R_1^l (k' \times n')$   
 $L_2^l (n' \times q')$ ,  $R_2^l (q' \times m')$   
 $S^{l-1} (k' \times q')$




# Parallel Implementation

## Distributed Binary Tree




- Distribution scheme at each level:
  - Calculate colors for sub-communicators
  - Calculate colors for each binary-tree nodes
  - Add nodes to MPI processes if node colors match sub-communicator colors
- Binary color calculation:
  - Colors (binary code) are unique for nodes/communicators in the same binary-tree level
  - Colors of child nodes/communicators are calculated from colors of their parents




## Schur-complement PCA Factorization

- What need to be stored in binary tree nodes?
  - All coupling matrices **C**
  - Diagonal block matrices of **A** at the lowest level
  - Block matrices of **A** at higher levels: store the **L**, **R**, and **S** matrices
- Block-based matrix distribution if a tree node is handled by more than one MPI rank
- Factorization: traversing the binary tree from the finest level up to a user-defined stop level while computing the **L**, **R**, and **S** matrices at each level using the **PCA** algorithm [3] (based on singular value decomposition - **SVD**)
- In-house distributed matrix-matrix (vector) multiplication (GEMM/GEMV)
- TPLs: singular value decomposition (SVD) and distributed linear equation solver GESV



Bottom-up Traversal for Factorization




Distributed matrix-matrix/vector multiplication  
GEMM/V

[3] [http://helper.ipam.ucla.edu/publications/setut/setut\\_7373.pdf](http://helper.ipam.ucla.edu/publications/setut/setut_7373.pdf)

## GCR Solver


- Generalized Conjugate Residual (GCR) method [4] was chosen for its simplicity and effectiveness
- Two most computational operations: applying preconditioner and matvec
- Applying preconditioner: using the bottom-up traversal as in the factorization and the in-house distributed GEMV with on-the-fly inverse matrices calculated from the stored **L**, **R**, and **S** matrices at each level
- Distributed binary-tree based matvec:
  - using the diagonal block matrices of **A** and the coupling matrices **C**
  - using the **entire** binary tree
  - output of a level becomes input for its upper level
  - using **MPI\_Allreduce** to gather final results




[4] Yousef Saad, Iterative Methods for Sparse Linear Systems; SIAM, 2003

## GCR Solver

- Generalized Conjugate Residual (GCR) method <sup>[4]</sup> was chosen for its simplicity and effectiveness
- Two most computational operations: applying preconditioner and matvec
- Applying preconditioner: using the bottom-up traversal as in the factorization and the in-house distributed GEMV with on-the-fly inverse matrices calculated from the stored **L**, **R**, and **S** matrices at each level
- Distributed binary-tree based matvec:
  - using the diagonal block matrices of **A** and the coupling matrices **C**
  - using the **entire** binary tree
  - output of a level becomes input for its upper level
  - using **MPI\_Allreduce** to gather final results




Bottom-up Traversal for Matvec



[4] Yousef Saad, Iterative Methods for Sparse Linear Systems; SIAM, 2003

## GCR Solver


- Generalized Conjugate Residual (GCR) method [4] was chosen for its simplicity and effectiveness
- Two most computational operations: applying preconditioner and matvec
- Applying preconditioner: using the bottom-up traversal as in the factorization and the in-house distributed GEMV with on-the-fly inverse matrices calculated from the stored **L**, **R**, and **S** matrices at each level
- Distributed binary-tree based matvec:
  - using the diagonal block matrices of **A** and the coupling matrices **C**
  - using the **entire** binary tree
  - output of a level becomes input for its upper level
  - using **MPI\_Allreduce** to gather final results




[4] Yousef Saad, Iterative Methods for Sparse Linear Systems; SIAM, 2003

## GCR Solver

- ❑ Generalized Conjugate Residual (GCR) method <sup>[4]</sup> was chosen for its simplicity and effectiveness
- ❑ Two most computational operations: applying preconditioner and matvec
- ❑ Applying preconditioner: using the bottom-up traversal as in the factorization and the in-house distributed GEMV with on-the-fly inverse matrices calculated from the stored **L**, **R**, and **S** matrices at each level
- ❑ Distributed binary-tree based matvec:
  - using the diagonal block matrices of **A** and the coupling matrices **C**
  - using the **entire** binary tree
  - output of a level becomes input for its upper level
  - using **MPI\_Allreduce** to gather final results




[4] Yousef Saad, Iterative Methods for Sparse Linear Systems; SIAM, 2003




# Numerical Results

## Simulation Setup and Convergence Behavior





- ❑ High-Q factor slotted cylindrical cavity at  $f = 1.1295\text{GHz}$  (near resonance) with 98,972 unknowns
- ❑ The matrix and RHS vector are generated by the method of moments code EIGER
- ❑ Schur-PCA tolerance:  $10\text{e-}5$  and GCR solver tolerance:  $10\text{e-}11$
- ❑ Test platform: LLNL's Lassen with POWER9 CPUs, V100 GPUs, gcc/8.3.1, cuda/11.8.0, essl/6.3.0.1, lapack/3.10.0-xl-2022.03.10, spectrum-mpi/rolling-release



- ❑ Lvl 1 → using the whole binary tree
- ❑ Lvl 8 (finest level) → block diagonal preconditioner
- ❑ The convergence rate is fastest with lvl 1 and reduces (i.e. more iterations) when we increase the stop level
  - The iteration change is significant from 7 (lvl 1) to 287 (lvl 2) and 1452 (lvl 8)



## CPU and GPU Executions at Different Stop Levels



- 16 computing nodes with 16 MPI ranks (1 rank per node): 44-core CPUs per rank vs. 1 GPU per rank
- The factorization time (blue) decreases as the stop level increases while the solve time (orange) increases with more iterations and becomes dominant in the total time (green)
- As the stop level increases, the memory requirement of the Schur-PCA factorization are less while the memory requirement of Krylov subspace is higher

- *Should use a stop level corresponding to the middle of the binary tree to balance between run time and memory footprint*
- GPU runs at stop levels of 3, 4, 5, 8 due to GPU memory limitation
- GPU execution is ~2.3x faster CPU execution

## Schur-PCA Preconditioner + GCR vs. ADELUS: Scalability



- The total CPU run times with stop levels 1 and 5 are compared with those of the direct solver ADELUS on 2 ranks, 4 ranks, 8 ranks, 16 ranks, 32 ranks
- Schur-PCA Preconditioner + GCR outperforms ADELUS
- Scalability needs optimization as the number of MPI ranks increases (current limitation: use of full-size vectors)

## Conclusions and Future Work

- Developed a parallel implementation of the Schur-complement PCA preconditioner on distributed-memory systems with Kokkos for performance portability
- Effective in solving ill-conditioned problems
- Schur-complement PCA preconditioner + GCR on CPUs is faster than ADELUS
- Future work:
  - Has just been integrated into Gemma: full performance with on-the-fly matrix construction will be done next
  - Performance optimization to achieve better scalability on larger computing systems for larger sized problems


**Thank You!**

# Backup

# Kokkos Kernels Overview

Kokkos Kernels is a library for *node-level*, performance-portable, computational kernels for sparse/dense linear algebra and graph operations, using the Kokkos

- KK is available publicly both as part of Trilinos and as part of the **Kokkos ecosystem** (<https://github.com/kokkos/kokkos-kernels>)
- **Building block** of a solver, linear algebra library that uses MPI and threads for parallelism, or it can be used stand-alone in an application
- Interfaces to **vendor-provided kernels** available in order to leverage their high-performance libraries



## Binary Tree-Based Inverse Matrix-Vector Multiplication



.....in compact form,

$$\begin{bmatrix} \mathbf{Y}'_1 \\ \mathbf{Y}'_2 \end{bmatrix} = \mathbf{B}_1^{-1} \begin{bmatrix} \mathbf{X}_1 \\ \mathbf{X}_2 \end{bmatrix} = \begin{bmatrix} (\mathbf{I} + \mathbf{L}_1 \mathbf{S} \mathbf{R}_1)(\mathbf{Y}_1 - \mathbf{L}_1 \mathbf{R}_1 \mathbf{Y}_2) \\ \mathbf{Y}_2 - \mathbf{L}_2 \mathbf{R}_2 (\mathbf{I} + \mathbf{L}_1 \mathbf{S} \mathbf{R}_1)(\mathbf{Y}_1 - \mathbf{L}_1 \mathbf{R}_1 \mathbf{Y}_2) \end{bmatrix}$$

where  $\mathbf{Y}_1 = \mathbf{A}_1^{-1} \mathbf{X}_1$  is an update on  $\mathbf{X}_1$  from child1

$\mathbf{Y}_2 = \mathbf{A}_2^{-1} \mathbf{X}_2$  is an update on  $\mathbf{X}_2$  from child2

$$\mathbf{S}^{k \times q} = (\mathbf{I} - \mathbf{R}_1 \mathbf{L}_2 \mathbf{R}_2 \mathbf{L}_1)^{-1} \mathbf{R}_1 \mathbf{L}_2$$

$\mathbf{M} \times \mathbf{V}$  procedure:

- 1) Declare 2 working vectors  $\mathbf{W}_1$  and  $\mathbf{W}_2$ , size  $\max(k, q)$
- 2) Compute  $\mathbf{W}_1^k = \mathbf{R}_1^{k \times n} \cdot \mathbf{Y}_1^n$
- 3) Update  $\mathbf{Y}_1^m = \mathbf{Y}_1^m - \mathbf{L}_1^{m \times k} \cdot \mathbf{W}_1^k$
- 4) Compute  $\mathbf{W}_1^q = \mathbf{R}_2^{q \times m} \cdot \mathbf{Y}_1^m$
- 5) Compute  $\mathbf{W}_2^k = \mathbf{S}^{k \times q} \cdot \mathbf{W}_1^q$
- 6) Update  $\mathbf{Y}_1^m = \mathbf{Y}_1^m + \mathbf{L}_1^{m \times k} \cdot \mathbf{W}_1^k$
- 7) Compute  $\mathbf{W}_1^q = \mathbf{R}_2^{q \times m} \cdot \mathbf{Y}_1^m$
- 8) Update  $\mathbf{Y}_2^n = \mathbf{Y}_2^n - \mathbf{L}_2^{n \times q} \cdot \mathbf{W}_1^q$