



Sandia
National
Laboratories

Exceptional service in the national interest

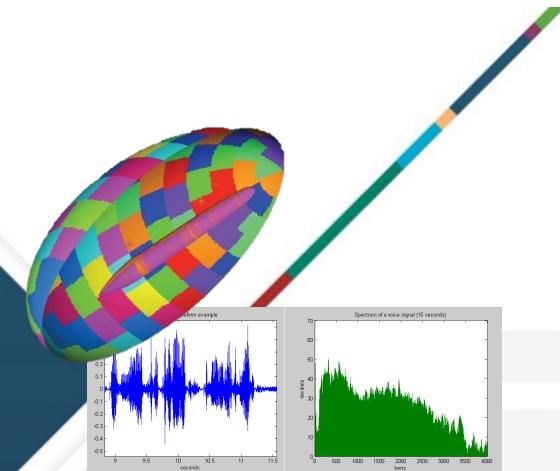
Performance of Modal Random Vibration Calculations for Structural Dynamics on Heterogeneous Computing Architectures

Julia Plews and Nathan Crane

Computational Solid Mechanics & Structural Dynamics
Sandia National Laboratories

Trilinos User-Developer Group Meeting 2023

October 30-November 2, 2023



COMP SIM
STRUCTURAL DYNAMICS

SAND2023-11591PE

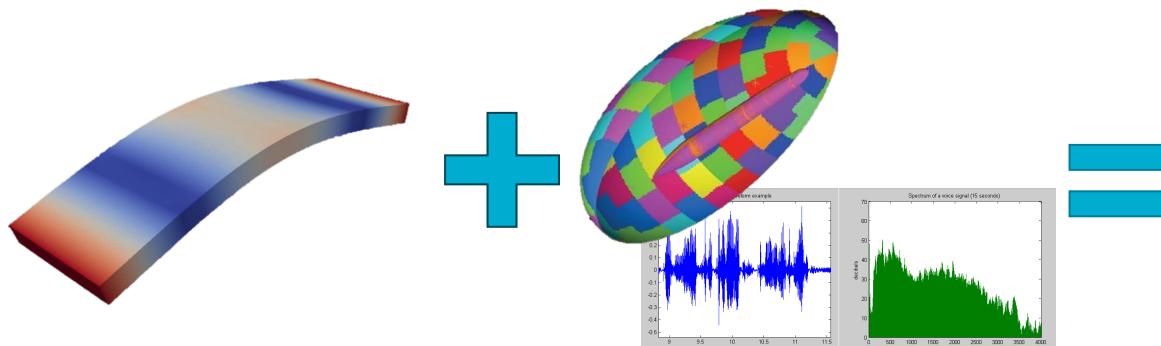
Sandia National Laboratories is a multimission laboratory managed and operated by National Technology and Engineering Solutions of Sandia LLC, a wholly owned subsidiary of Honeywell International Inc. for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA0003525.





What is Modal Random Vibration?

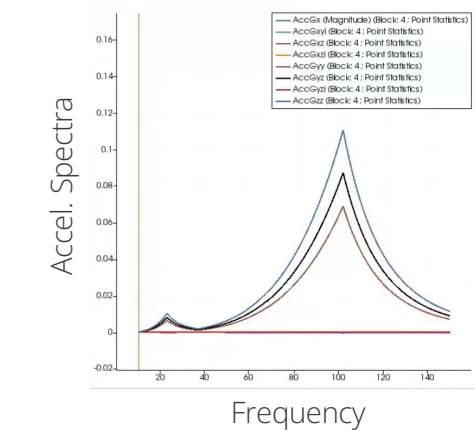
Modal random vibration is a means to evaluate the statistical behavior of a structure in a steady-state vibration environment



Natural frequency
mode shapes

Statistical description of
input loads

- Power spectra in frequency domain
- Correlation between input loads



Statistical **dynamic system**
behavior

- Acceleration
- Displacement
- Stress

Image: https://en.wikipedia.org/wiki/Spectral_density

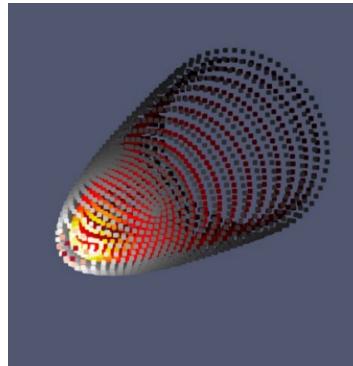


What is Modal Random Vibration?

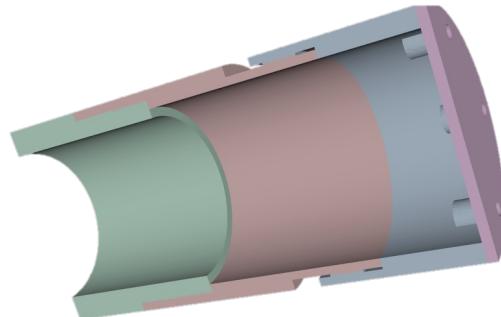
An efficient alternative to performing **large ensembles of transient structural analyses** to understand response of structures under random excitation in environments such as



Transportation



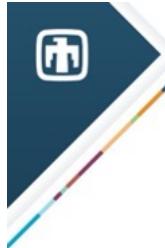
Flight and atmospheric reentry



Test fixtures and experimental configurations

Typical probabilistic quantities of interest include

- margin to material yield,
- acceleration spectra near or inside a structural component,
- fatigue crack initiation and growth, etc.



Modal Random Vibration mathematics

$$S(\omega) = Z(\omega)_{\text{modes} \times \text{loads}} E(\omega)_{\text{loads} \times \text{loads}} Z^T(\omega)_{\text{loads} \times \text{modes}}$$

$$A(\omega) = H^*(\omega)_{\text{modes} \times \text{modes}} S(\omega)_{\text{modes} \times \text{modes}} H(\omega)_{\text{modes}}$$

ω = response frequency

$A(\omega)$ = response output (maybe complex-valued)

$H(\omega)$ = transfer function (complex-valued)

$S(\omega)$ = modal input power spectrum (complex-valued)

$Z(\omega)$ = modal excitation for load (real-valued)

$E(\omega)$ = input load magnitudes and correlations (complex-valued, input)

Typically engineers deal with up to thousands each of modes, loads, frequencies, and output locations, yielding possibly *trillions of computations* for a single simulation

An opportunity to utilize novel high-performance computing platforms to drastically reduce simulation turnaround times

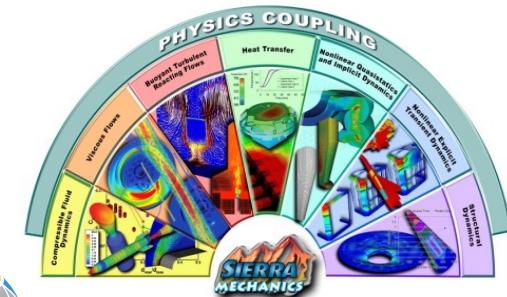


Sierra Structural Dynamics simulations at Sandia

Sierra/SD (Structural Dynamics):
massively parallel C++ finite element
analysis code in the **Sierra** suite of tools

- In development for 25+ years
- Designed to run on the fastest supercomputers in the world featuring **CPU and GPU acceleration**
- Support national security missions

SD simulations are used to predict the behavior of a wide variety of systems, components, and experimental configurations

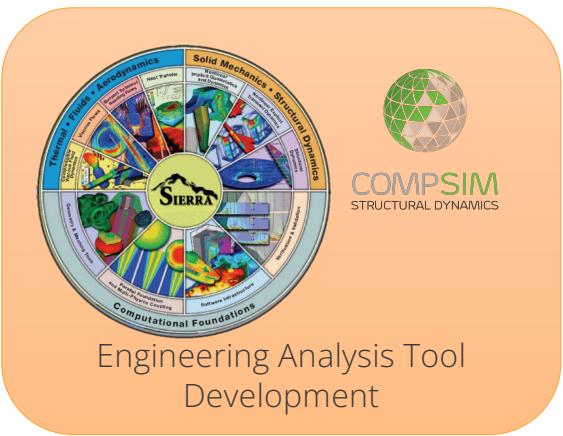


COMP SIM
SOLID MECHANICS





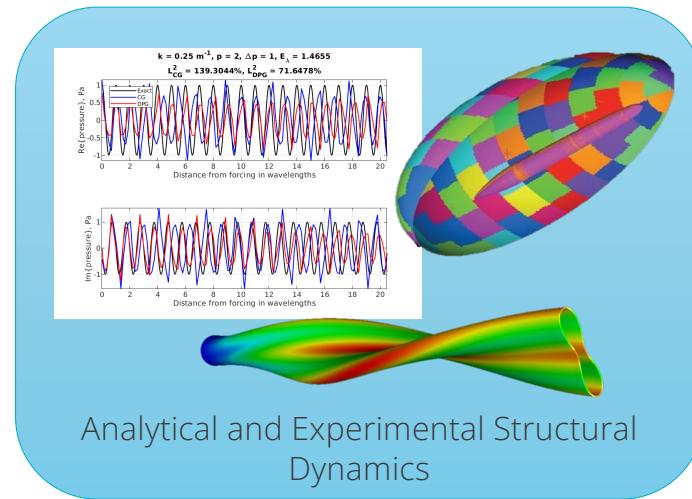
Computational science and engineering research for success in large SD simulations



Engineering Sciences



Computing Research





Porting Sierra/SD to GPUs: Kokkos performance portability library

Sandia-developed C++ Performance Portability Programming EcoSystem¹

- Write algorithms just once, execute on any architecture

```
Kokkos::View<double*, Kokkos::CudaSpace>
    v("myData", 1024);
auto vHost = Kokkos::create_mirror(v);
... // do work on vHost
Kokkos::deep_copy(v, vHost);
```

- Automatically move data between memory spaces (e.g., CUDA to host and back)

```
KokkosBlas::gemm(Sal::ExecSpace(), "N", "N", 1.0, z, Eij,
                  0.0, Z_times_E);
KokkosBlas::gemm(Sal::ExecSpace(), "N", "C", 1.0, ZtimesE,
                  z, 0.0, Sij);
```

- Provide an agnostic interface to vendor-supplied linear algebra operations (e.g., cuBlas)

```
struct MyKernel
{
    void operator()(const size_t i) const {
        printf("Hello, world!\n");
        myGpuData(i) = 2 * i;
    }

    Kokkos::View<double*> myGpuData;
};

Kokkos::parallel_for(1024, MyKernel());
```

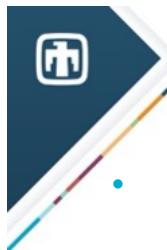
- Execute kernels in parallel through a tailored template interface

Kokkos on Github:
<https://github.com/kokkos>



1. [Kokkos: Enabling manycore performance portability through polymorphic memory access patterns](#)

H. C. Edwards , C. R. Trott , D. Sunderland, Kokkos, J. Parallel and Distributed Computing, v.74 n.12, p.3202-3216, Dec. 2014



Historical performance of Modal Random Vibration in Sierra/SD

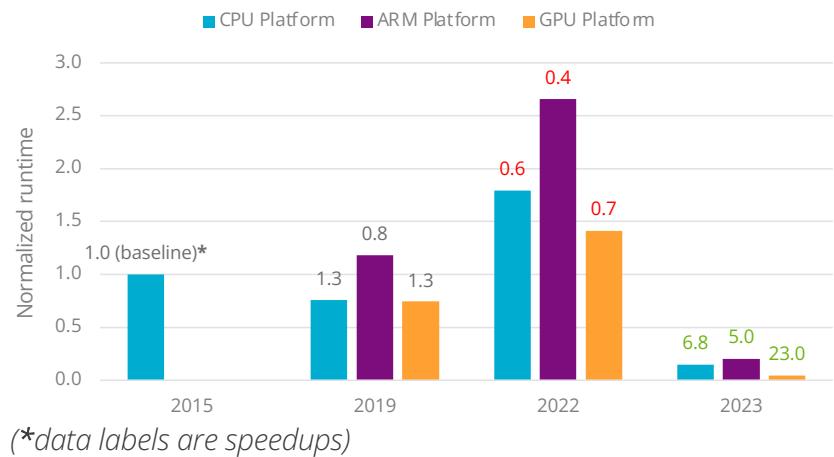
- Exemplar model performance:
Medium-fidelity simulation of vibration of a system in atmospheric reentry environment



ARM platform: "Astra," 28 cores x 2 sockets Cavium Thunder-X2 per node



GPU platform: "Sierra," 44 cores IBM Power9 + 4x NVIDIA V100 per node



- 2015: Historical performance, CPU only
- 2019: Initial Kokkos port to both ARM and GPU
- 2022: First attempt at GPU optimization (lessons learned!)
- 2023: Second, successful attempt at GPU optimization based on KokkosKernels



Historical performance of Modal Random Vibration in Sierra/SD

2019: Initial port to GPU and ARM based on Kokkos

Focused on a single performance hotspot: nodal output at each frequency

- Overall complexity $O(\text{frequencies} * \text{nodes} * \text{modes}^2)$
- GPU parallelization *only* over nodes; $O(1000)$ concurrency insufficient to saturate GPU
- Modal input power spectrum (S) calculation *not* ported to GPU

```
for (int i = 0; i < numFreq; ++i) {
    Kokkos::parallel_for("NodalOutput", numNodes,
        KOKKOS_LAMBDA(int j) {
            outputVal(j) = Calc_H_S_Ht(H(j), S);
        });
}
```

Serialized calculation per node

No observable performance benefit from GPU use: small concurrency and too many serialized computations



Historical performance of Modal Random Vibration in Sierra/SD

2021: (Unsuccessful) first attempt at GPU optimization

GPU optimization focused on calculation of matrix S (input power spectrum):

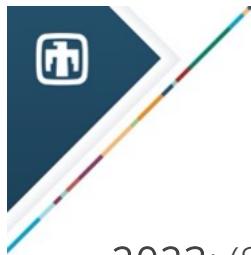
- $S = Z_{\text{modes} \times \text{loads}} E_{\text{loads} \times \text{loads}} Z^T_{\text{loads} \times \text{modes}}$
- Kokkos “**multi-dimensional range**” to consolidate 4 nested loops

```
Kokkos::MDRangePolicy<Kokkos::Rank<4>> range({0, 0, 0, 0}, {nModes, nModes, nLoads, nLoads});  
Kokkos::parallel_for(  
    "Scalc", range, KOKKOS_LAMBDA(int iMode, int jMode, int iLoad, int jLoad) {  
        unsigned iModeIdx = modeIndices(iMode);  
        unsigned jModeIdx = modeIndices(jMode);  
        Kokkos::atomic_add(&S(iModeIdx, jModeIdx),  
                           Z(iModeIdx, iLoad) * E(iLoad, jLoad) * Z(jModeIdx, jLoad));  
    } );
```

Accessing data at non-contiguous locations

$O(\text{modes}^2 * \text{loads}^2)$

Performance actually *degraded* due to atomics, frequent data access indirection, and high operator complexity



Historical performance of Modal Random Vibration in Sierra/SD

2023: (Successful) second attempt at GPU optimization

- $S = Z_{\text{modes} \times \text{loads}} E_{\text{loads} \times \text{loads}} Z^T_{\text{loads} \times \text{modes}}$
- Step one: **remove modal data access indirection**

```
RemoveInactiveModes(S);
RemoveInactiveModes(Z);
Kokkos::MDRangePolicy<Kokkos::Rank<4>> range({0, 0, 0, 0}, {nModes, nModes, nLoads, nLoads});
Kokkos::parallel_for(
    "Scalc", range, KOKKOS_LAMBDA(int iMode, int jMode, int iLoad, int jLoad) {
        Kokkos::atomic_add(&S(iMode, jMode),
                          Z(iMode, iLoad) * E(iLoad, jLoad) * Z(jMode, jLoad));
    });
}
```

Data access is now contiguous, but
complexity is still $O(\text{modes}^2 * \text{loads}^2)$



Historical performance of Modal Random Vibration in Sierra/SD

2023: (Successful) second attempt at GPU optimization

- $ZE = Z_{\text{modes}} \times \text{loads } E_{\text{loads}} \times \text{loads}$
- $S = ZE_{\text{modes}} \times \text{loads } Z^T_{\text{loads}} \times \text{modes}$
- Step two: introduce temporary product, use performance-portable BLAS provided by KokkosKernels (optimized dense linear algebra)

```
RemoveInactiveModes(S);  
RemoveInactiveModes(Z);  
KokkosBlas::gemm(ExecSpace(), "N", "N", 1.0, Z, E, 0.0, ZtimesE);  
KokkosBlas::gemm(ExecSpace(), "N", "C", 1.0, ZtimesE, Z, 0.0, S);
```

Cleaner code, lower $O(\text{modes}^2 * \text{loads})$ complexity, and truly performance-portable



Summary and impacts on statistical analysis of structures at Sandia

	Spring 2022 baseline	Fall 2022	Spring 2023
CPU	--	~3x slower than Spring 22	~4x faster than Spring 22
GPU	on par with CPU	~1.2x slower than Spring 22 on CPU	~12x faster than Spring 22 on CPU

Performance changes in actual SD analyst model, relative to Spring 2022 version of Sierra/SD on CPU HPC platform

Through a combination of

- advances in GPU architectures,
- convenient code porting utilities,
- and algorithm optimizations,



modal random vibration computations in Sierra/SD that *used to take days* may **now take minutes**, enabling statistical analyses of structural designs in real time